

Figure 4 - Learned policy vs closed-form optimum ($n = 5$ nodes = 4 segments, $p_{\text{gen}} = 0.4$, $p_{\text{swap}} = 0.7$)

